

Week 9 Exercises: State Machines and Markov Decision Processes

🔖 Bookmark this page

Week 9 Exercises: State Machines and Markov Decision Processes

(1.0 points possible)

For these exercises, you should read the notes on [State Machines and Markov Decision Processes](#).


1) State Machines

State machines are a powerful and pervasive concept in information processing, playing important roles in computer science, signal processing, and control. Fundamentally, you can think of a state machine as a module that maps a sequence of input values  $x_1, \dots, x_T$  (they could be numbers, words, graphs, or anything) into a sequence of output values (also of any type)  $y_1, \dots, y_T$ .

In the simplest case, we might have a simple functional dependence, so that  $y_t = h(x_t)$ , but more generally, we are interested in the case where  $y_t$  might potentially depend on all previous input values, so  $y_t = h(x_1, \dots, x_t)$ . Generally, describing a functional dependency like this is very difficult, because the function  $h$  seems to need a different number of arguments on each time step. However, we can describe a very large class of  $h$  functions compactly, as a recursive combination of two functions,  $f$  and  $g$ , with an intermediate *state* value capturing the information that flows between applications of these functions. Mathematically, we have

$$\begin{aligned} y_t &= g(s_t) \\ s_t &= f(s_{t-1}, x_t) \\ s_0 &= 0 \end{aligned}$$

Here is a figure illustrating the process: you can think of it as a circuit with the state fed back into  $f$ , but with a one-step delay.



If  $f$  and  $g$  are linear functions, then this is an LTI (linear time-invariant) system, a popular model in control systems analysis. In general, an LTI system can depend on a finite number of previous input values and a finite number of previous output values; we can model that dependence by allowing  $s_t$  to be a vector, storing as many of these values as we'd like.

If  $x_t$ ,  $y_t$ , and  $s_t$  are all elements of finite sets, then this describes an FSM (finite-state machine), a popular model in computer science and some kinds of discrete control.

For each of the following state machines, provide the output sequence  $[y_1, y_2, \dots, y_T]$  given the input sequence  $[x_1, x_2, \dots, x_T]$ :

1A)

```
s_0 = 0
f(s, x_i) = max(s, x_i)
g(s) = s * 2
Input: [0, 1, 2, 1]
```

Enter a Python list of four numbers: [0, 2, 4, 4]

Submit View Answer 100.00%

You have infinitely many submissions remaining.

1B)

```
s_0 = (0, 0)
f(s, x_i) = (s[0] + x_i, s[1] + 1)
g(s) = s[0] / s[1]
Input: [0, 1, 2, 1]
```

Enter a Python list of four numbers: [0, 0.5, 1, 1]

Submit View Answer 100.00%


You have infinitely many submissions remaining.

2) Tiny Policy Evaluation


Consider an MDP with states (0, 1, 2, 3) and actions ('b', 'c'). The reward function is:

$$R(s, a) = \begin{cases} 1 & \text{if } s = 1 \\ 2 & \text{if } s = 3 \\ 0 & \text{otherwise} \end{cases}$$

You get the reward associated with a state on the step when you exit that state. The transition function for each action is below, where  $T[i, x, j]$  is the conditional probability  $P(s_{t+1} = j | a = x, s_t = i)$ . The state transition diagrams are given as well, but you may find it easier to calculate the state values using the transition matrices.



$$T(s_t, 'b', s_{t+1}) = \begin{bmatrix} 0.0 & 0.9 & 0.1 & 0.0 \\ 0.9 & 0.1 & 0.0 & 0.0 \\ 0.0 & 0.0 & 0.1 & 0.9 \\ 0.9 & 0.0 & 0.0 & 0.1 \end{bmatrix}$$



$$T(s_t, 'c', s_{t+1}) = \begin{bmatrix} 0.0 & 0.1 & 0.9 & 0.0 \\ 0.9 & 0.1 & 0.0 & 0.0 \\ 0.0 & 0.0 & 0.1 & 0.9 \\ 0.9 & 0.0 & 0.0 & 0.1 \end{bmatrix}$$

Note that the **only** effect of the action is to change the transition probability from state 0 (the first row of the transition matrix: rows correspond to the input states, and columns correspond to the output states).

We would like to find the value function associated with the policy that **always** chooses action 'c'.

2A) What are the horizon 0 undiscounted values of the states under this policy?

Enter a Python list of four numbers: [0, 0, 0, 0]

Submit View Answer 100.00%

You have infinitely many submissions remaining.

2B) For horizon 1?

Enter a Python list of four numbers: [0, 1, 0, 2]

Submit View Answer 100.00%

You have infinitely many submissions remaining.

2C) For horizon 2?


Enter a Python list of four numbers: [0.1, 1.1, 1.8, 2.2]

Submit View Answer 100.00%

You have infinitely many submissions remaining.

This page was last updated on Saturday January 18, 2020 at 01:16:11 PM (revision 043d9ae).

Sec None Release 1900-01-01 00:00:00 Due 9999-12-31 23:59:59



Powered by CAT-SOOP v14.0.0.dev144.

CAT-SOOP is free/libre software, available under the terms

of the GNU Affero General Public License, version 3.

[Download Source Code](#)

[Javascript License Information](#)